

Deployment of a Docker container of “Yolo” in two cloud services: as an Azure container instance and a virtual machine

José Luis Puente Bodoque

June 25, 2023

Abstract

In this report I comment the steps followed in order to deploy and test a Yolo app project as mandatory assignment in module 13: Cloud Computing.

Contents

1	Testing the app locally	2
2	Building the container	2
3	Testing the container locally	3
4	Deployment in Azure as an Azure container instance	3
4.1	Pushing my container to hub.docker.com	3
4.2	Deployment and test of an Azure Container Instance	5
4.3	Deployment and test in a virtual machine	5

1 Testing the app locally

To test the app locally, I open a terminal window and enter the following command to run the app.

```
python app.py
```

Once the app is running I can test its functionality. I can access it at:

```
localhost:4000/upload
```

For some Real Madrid starter line-up image in JPEG format, the result is the one below.



Figure 1: Testing the app locally

2 Building the container

The project already provides a Dockerfile in the same directory as the app.py so I just run these Docker commands:

```
# Set the default Docker context
docker context use default
```

```
# Build the container  
docker build -t yolo .
```

I received the error “failed to fetch from deb.debian.org...” when I entered the second command. I solved it¹ by entering the following command to clear the Docker cache:

```
docker system prune --all
```

3 Testing the container locally

To verify that the container is working I entered the command:

```
# Verify that the container is working  
docker run -p 4000:4000 yolo
```

Then I tested it at:

```
http://127.0.0.1:4000/upload
```

4 Deployment in Azure as an Azure container instance

I have deployed a container from one already pushed to hub.docker.com. So the first step is to create an account in hub.docker.com and then create a repository.

4.1 Pushing my container to hub.docker.com

My public repository for yolo project is the one below.

In order to push my container to my repository and then test it locally, I opened a window terminal in the same directory as the Yolo app project and entered:

```
# Tag my image:  
docker tag yolo joseluispuente/yolo:some-version  
  
# Upload it to hub.docker.com:
```

¹**This** answer truly deserves all the merit.

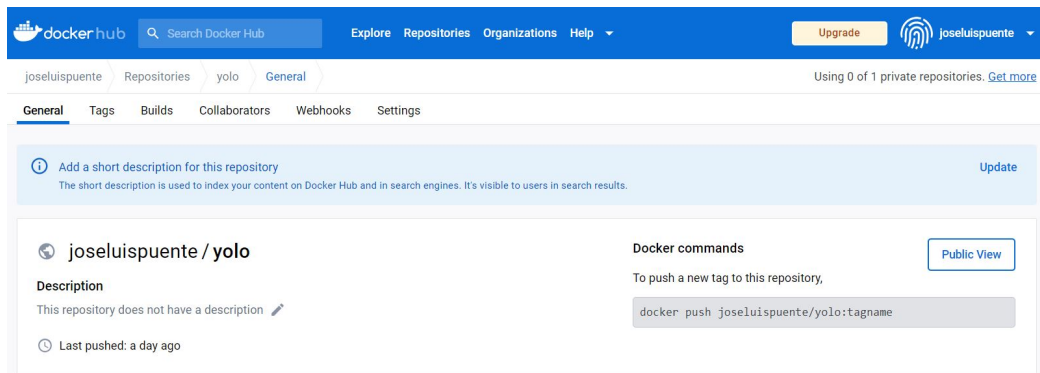


Figure 2: My repository at hub.docker.com

```
docker push joseluispuente/yolo:some-version

# Run it just by typing (nameless container):
docker run -p 4000:4000 joseluispuente/yolo:some-version

# Test it at:
localhost:4000/upload
```

I can see that my containers are built and visible on Docker Desktop.

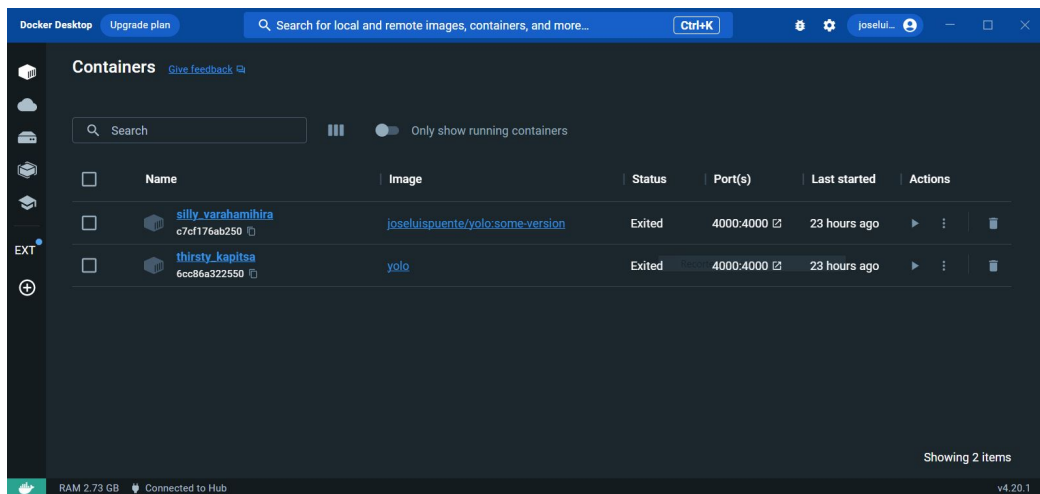


Figure 3: My containers on Docker Desktop

4.2 Deployment and test of an Azure Container Instance

I have created a new Azure container instance by following the steps indicated in class slides. Then I click on Run and the resulting window is shown below. Note the public IP is 20.13.117.190.

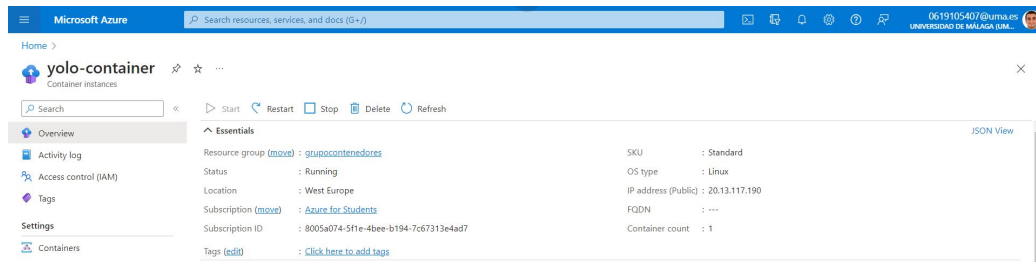


Figure 4: My Azure container instance of Yolo

Finally I may test the running container by using its public IP and port number 4000.



Figure 5: Test of y Azure container instance of Yolo

4.3 Deployment and test in a virtual machine

First step is to create a new virtual machine as resource in Azure.

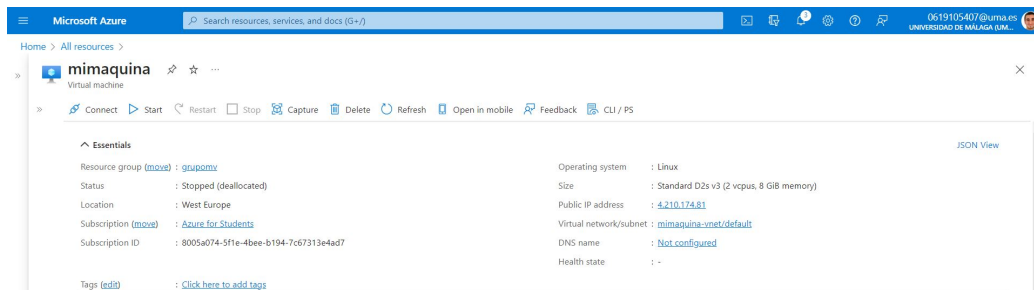


Figure 6: My virtual machine

Next step is to connect to it using PuTTY with the user `userazure` and public IP 4.210.174.81, so that the host name to type in the logging tap is `userazure@4.210.174.81`. The private key must suit PuTTY format, that is, ppk (PuTTY Private Key) format. As the key I downloaded when I created my virtual machine has pem format, I used PuTTYgen to convert it to ppk format.

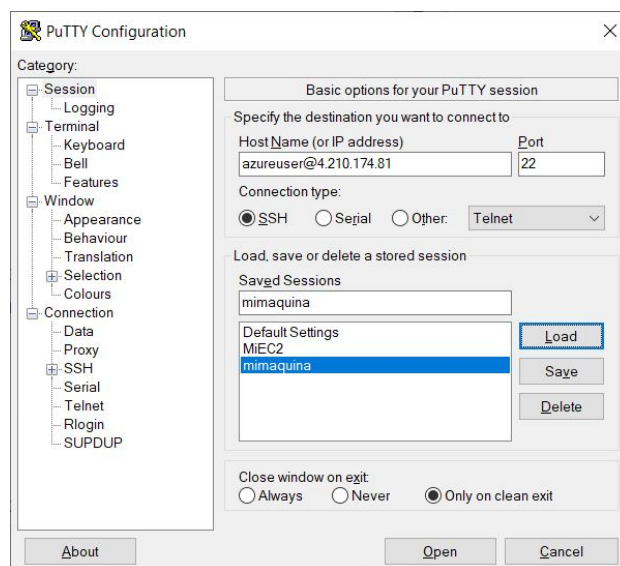


Figure 7: Putty login

Once I was already connected I installed Python at first.

```
sudo apt-get install python
```

Then I tried to create a virtual environment, but a helpful error message suggested me to install `python3-venv` before.

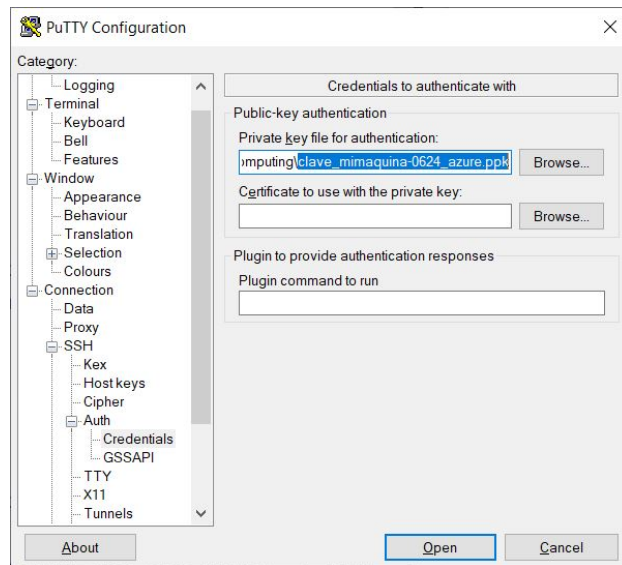


Figure 8: My PuTTY private key imported

```
sudo apt-get install python3-venv
```

Now I could create a virtual environment (.venv):

```
python3 -m venv .venv
```

And I activated it:

```
source .venv/bin/activate
```

Now I installed all dependencies specified in requirements.txt:

```
python -m pip install -r requirements.txt
```

And when I ran the app.py I received an error related to the cv2 package. I found the solution in this way:

```
sudo apt-get update
sudo apt-get install ffmpeg libsm6 libxext6 -y
```

These two last commands install the cv2 dependencies that aren't present just by installing the opencv-python library.

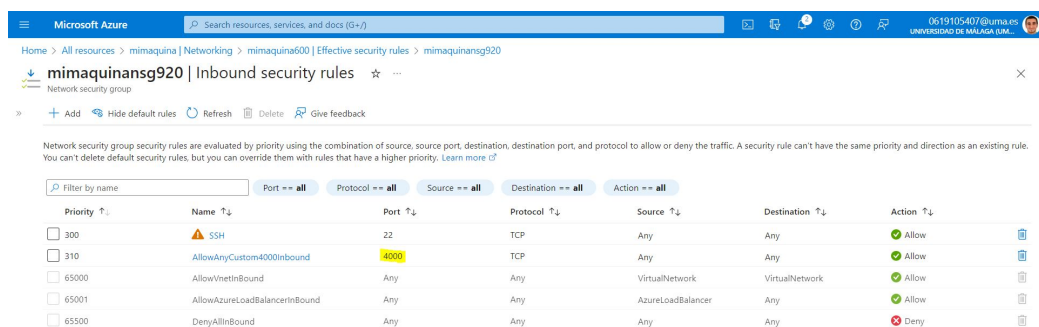
Now I could successfully launch the app.py.

In order to enter from the address bar of the browser it was necessary for me to open the port 4000 in the virtual machine. To do that, I went to

```
azureuser@mimaquina: ~/yolo_base
Setting up vdpau-driver-all:amd64 (1.4-3) ...
Setting up libavfilter7:amd64 (7:4.3.6-0+deb11u1) ...
Setting up libavdevice58:amd64 (7:4.3.6-0+deb11u1) ...
Setting up ffmpeg (7:4.3.6-0+deb11u1) ...
Processing triggers for man-db (2.9.4-2) ...
Processing triggers for libc-bin (2.31-13+deb11u6) ...
Processing triggers for libgdk-pixbuf-2.0-0:amd64 (2.42.2+dfsg-1+deb11u1) ...
(.venv) azureuser@mimaquina:~/yolo_base$ python app.py
[INFO] loading YOLO from disk...
* Serving Flask app 'app' (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on all addresses (0.0.0.0)
  WARNING: This is a development server. Do not use it in a production deployment.
* Running on http://127.0.0.1:4000
* Running on http://10.0.0.4:4000 (Press CTRL+C to quit)
178.239.240.126 - - [25/Jun/2023 19:46:50] "GET / HTTP/1.1" 404 -
178.239.240.126 - - [25/Jun/2023 19:46:50] "GET /favicon.ico HTTP/1.1" 404 -
178.239.240.126 - - [25/Jun/2023 19:46:57] "GET /upload HTTP/1.1" 200 -
```

Figure 9: Running app.py

my virtual machine overview, then Networking and add a new inbound port rule.



Priority	Name	Port	Protocol	Source	Destination	Action
300	SSH	22	TCP	Any	Any	Allow
310	AllowAnyCustom4000Inbound	4000	TCP	Any	Any	Allow
65000	AllowVnetInBound	Any	Any	VirtualNetwork	VirtualNetwork	Allow
65001	AllowAzureLoadBalancerInBound	Any	Any	AzureLoadBalancer	Any	Allow
65500	DenyAllInBound	Any	Any	Any	Any	Deny

Figure 10: Inbound port rule

This allowed me to enter the application.

Finally I test it with the same image. Note that the public IP is the virtual machine, not the container instance one.



Figure 11: Access to app running on a virtual machine from a browser